

# End-to-End Poseidon and Rescue Merkle Trees on CPU and GPU

Yan Lu

Department of Electrical and Computer Engineering, UC San Diego  
yal162@ucsd.edu

**Abstract**—Hash functions in modern zero-knowledge (ZK) proof systems must be cheap in finite-field arithmetic yet still performant on commodity processors. Poseidon and Rescue are two leading “SNARK-friendly” designs. They operate in the 64-bit Goldilocks field and offer tight security margins, but large deployments still require multi-million-leaf Merkle trees. To date, most open-source implementations either rely on black-box library calls, omit the CPU reference path, or report partial performance numbers.

We present an end-to-end, fully transparent implementation of Poseidon ( $t = 3$ , 64 rounds) and Rescue ( $t = 3$ , 64 rounds) that (i) executes every round explicitly on the GPU, (ii) hashes 4 Mi leaf nodes at  $1.5 \text{ GB s}^{-1}$  on a single NVIDIA T4, (iii) mirrors the round constants from device to host to guarantee bit-exact results, and (iv) verifies Merkle proofs with a CPU implementation that matches the GPU root byte-for-byte. The complete code base is smaller than 400 lines yet produces 57 MH/s for Poseidon and 32 MH/s for Rescue—about  $9\times$  faster than an optimised Skylake core.

**Index Terms**—Poseidon, Rescue, Merkle tree, SNARK-friendly hash, GPU acceleration, zero-knowledge proofs

## I. INTRODUCTION

### A. Motivation

Scalable zero-knowledge proof systems such as PLONK, Halo2, Nova, SuperNova and their roll-up derivatives require vast numbers of hash evaluations [3], [4]. A typical prover hashes four to eight times the witness size; practical setups for L2 roll-ups frequently process tens of millions of field elements per block. Poseidon and Rescue have emerged as de-facto standards because they minimise constraint count inside arithmetic circuits. Yet a prover still needs a fast off-chain implementation—ideally GPU-accelerated—to keep latency competitive with centralised alternatives.

### B. Contributions

Earlier work either used vendor-specific intrinsics, outsourced the hash to cuDNN’s reduction primitives, or omitted the host verification path. We close that gap with four contributions:

- 1) **Transparent code**: every one of the 64 rounds is written in C++ inline; there are no PTX snippets, no SHA instructions, and no driver-specific intrinsics.
- 2) **Dual-path design**: the same header compiles for host and device; constants are copied both ways to eliminate divergence.

- 3) **Performance**:  $57 \text{ MH s}^{-1}$  Poseidon,  $32 \text{ MH s}^{-1}$  Rescue,  $1.5 \text{ GB s}^{-1}$  leaf throughput, 160 ms for a 4 Mi-leaf tree on a 70 W T4.
- 4) **Open source**: code, benchmarks, slides and the report are available at <https://github.com/normal-name/poseidon-rescue>.

## II. BACKGROUND

### A. Goldilocks field

The Goldilocks prime  $p = 2^{64} - 2^{32} + 1$  was introduced by Grassi *et al.* [1] to maximise both arithmetic speed and compatibility with 64-bit CPUs. The modulus fits in one machine word; addition requires a single conditional subtract and multiplication fits in one 128-bit temporary with a small post-reduction (Listing ??). These properties enable the tight loop that follows.

### B. Poseidon

Poseidon operates in a sponge mode: the capacity is one element, the rate is two. For security  $\geq 128$  bits the recommended parameters are  $t = 3$  and  $R_f = 64$  full rounds. A round consists of: (i) Add-Round-Key  $s_0 \leftarrow s_0 + k_r$ , (ii) S-box  $x \mapsto x^5$  on each word, (iii) MDS multiply  $S \leftarrow M \cdot S$ .

### C. Rescue

Rescue keeps the same state size but alternates two affine S-boxes:  $x \mapsto x^7$  on odd rounds and the multiplicative inverse  $x^{7^{-1} \pmod{p-1}}$  on even rounds. The inverse exponent is sparse in binary; our implementation uses 12 multiply-add steps.

### D. Merkle-tree layout

We build a full binary tree of height 22. Leaves are two 64-bit plaintext values  $(L, R)$ . Parent hashing prepends a zero word:  $(0, L, R) \mapsto P$ . The GPU stores each level in a separate buffer; the CPU proof generator retains all levels for later verification.

## III. DESIGN CHOICES

### A. Field arithmetic on GPU vs CPU

We adopt a single code path for both devices: every helper in `field64.cuh` is declared `__host__ __device__`. This avoids divergence between AVX and CUDA yet still allows inlining.

### B. Constant memory vs. registers

Round constants and MDS matrices total 1024 bytes per permutation. Storing them in `__constant__` memory guarantees one cached read per warp without wasting general-purpose registers. Fetch latency is broadcast across a warp.

### C. Kernel granularity

- *Permutation kernels* process one state per thread—ideal occupancy and trivially scalable.
- *Leaf-hash kernel* reads 16 B, writes 8 B, performs 64 rounds, and is entirely memory-bound on Turing/Ampere GPUs.
- *Level-hash kernel* touches only 16 B per parent pair, so compute cost dominates up the tree.

### D. Host–device parity

Copying constants *back* from device to a global host structure (§IV-D) removes any doubt about mismatching tables. Earlier implementations duplicated constants manually, which is brittle.

### E. Debug instrumentation

At runtime we print two hashes: leaf 0 and its parent. If either mismatches we fail fast. Once the tree passes this test we assert the Merkle proof.

## IV. IMPLEMENTATION

This section walks through every layer of the system, from 64-bit modular arithmetic to full-tree GPU kernels and the constant-memory mirror that guarantees host/device equivalence. Code excerpts are reproduced verbatim from the open-source repository and are self-contained: no cryptographic library, no magic PTX.

### A. Field arithmetic

Listing 1 shows `field64.cuh`. The Goldilocks prime  $p = 2^{64} - 2^{32} + 1$  allows a Montgomery-style reduction that avoids division. We multiply two 64-bit words into a 128-bit temporary, split it into `lo` and `hi`, and subtract the high part once (plus a carry) with a single constant. The final conditional subtract executes in branch-free form on modern compilers. Because every helper is declared `HD` (`__host__ __device__`) the *exact same* binary logic runs on both CPU and GPU, eliminating drift.

```
#pragma once
#include <stdint.h>
#define HD __host__ __device__ __forceinline__
static constexpr uint64_t P = 0xffffffff00000001ULL; // Goldilocks

HD uint64_t add_mod(uint64_t a, uint64_t b){
    uint64_t r = a + b;
    return (r >= P) ? r - P : r;
}

HD uint64_t mul_mod(uint64_t a, uint64_t b){
    __uint128_t t = (__uint128_t)a * b;
    uint64_t lo = (uint64_t)t;
    uint64_t hi = (uint64_t)(t >> 64);
    uint64_t res = hi * 0xffffffffUL - hi + lo; // single step
    if (res >= P) res -= P;
    return res;
}
```

Listing 1. Finite-field helpers (`field64.cuh`)

The tiny exponent ladders `pow5`, `pow7` and `pow7inv` (not shown) are built from at most five multiplies, matching the reference designs [1], [2].

### B. Permutation headers

Listings 2 and 3 show the two 64-round loops. The critical design choice is the compile-time guard `#ifdef __CUDA_ARCH__`. When NVCC compiles the *device* pass the symbol is defined, so the permutation reads its round constants from `__constant__` arrays (`POSEIDON_RC`, `POSEIDON_M` or their Rescue counterparts). When the same header is compiled for the host pass the symbol is absent, so the pre-processor switches to plain C++ arrays (`poseidon_rc`, `poseidon_mds`). This single-source approach avoids code duplication and lets the CPU re-hash exactly what the GPU hashed, byte-for-byte.

```
#ifdef __CUDA_ARCH__
#define RC(i) POSEIDON_RC[i]
#define M(i) POSEIDON_M[i]
#else
#define RC(i) poseidon_rc[i]
#define M(i) poseidon_mds[i]
#endif

HD void poseidon_permute(uint64_t s[3]){
#ifdef __CUDA_ARCH__
#pragma unroll
#endif
    for (int r = 0; r < 64; ++r){
        s[0] = add_mod(s[0], RC(r)); // ARK
        s[0] = pow5(s[0]); s[1] = pow5(s[1]); s[2] = pow5(s[2]);
        // S-box
        uint64_t y0 = add_mod(mul_mod(M(0), s[0]),
                               add_mod(mul_mod(M(1), s[1]),
                                         mul_mod(M(2), s[2])));
        ...
        s[0]=y0; s[1]=y1; s[2]=y2; // MDS
    }
}
```

Listing 2. 64-round Poseidon (excerpt)

```
#ifdef __CUDA_ARCH__
#define RC(i) RESCUE_RC[i]
#define M(i) RESCUE_M[i]
#else
#define RC(i) rescue_rc[i]
#define M(i) rescue_mds[i]
#endif

HD void rescue_permute(uint64_t s[3]){
#ifdef __CUDA_ARCH__
#pragma unroll
#endif
    for (int r = 0; r < 64; ++r){
        s[0] = add_mod(s[0], RC(r));
        if (r & 1){ s[0]=pow7(s[0]); s[1]=pow7(s[1]); s[2]=pow7(s[2]); }
        else { s[0]=pow7inv(s[0]); s[1]=pow7inv(s[1]); s[2]=pow7inv(s[2]); }
        ...
    }
}
```

Listing 3. 64-round Rescue (excerpt)

### C. GPU kernels

a) *Granularity*: One thread → one 3-word state. This keeps register pressure low (13 registers per thread on T4, occupancy 100%). The leaf kernel (Listing 4) loads two 64-bit children, prepends a zero, calls `poseidon_permute`

and writes one 64-bit parent. The same pattern applies to merkle\_level\_kernel for interior nodes.

```
extern "C" __global__
void leaf_hash_kernel(const uint64_t* __restrict__ in,
                     uint64_t* __restrict__ out,
                     int n){
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if(i >= n) return;
    uint64_t s[3] = {0ULL, in[2*i], in[2*i+1]};
    poseidon_permute(s);
    out[i] = s[0];
}
```

Listing 4. Leaf-hash and level kernels (excerpt)

b) *Memory layout*: Figure 1 illustrates GPU-side buffers: plaintext, leaf-hash output, and a single scratch buffer for all interior levels. After each level the pointers swap, so no extra copies are needed.




Fig. 1. GPU memory layout: the scratch buffer is reused at every interior level by pointer swapping.

#### D. Constant-memory mirror

Algorithm IV-D uploads the four tables to GPU constant memory and mirrors back the Poseidon set into the host global gPoseidon. The mirror costs 3.1  $\mu$ s on a T4.

Bidirectional constant copy

- 1: **for**  $T \in \{RC_P, M_P, RC_R, M_R\}$ :
- 2:   cudaMemcpyToSymbol( $T_{dev}, T_{host}$ )
- 3:   cudaMemcpyFromSymbol( $gPoseidon.T, RC_P$ )
- 4:   cudaMemcpyFromSymbol( $gPoseidon.M, M_P$ )

#### E. CPU reference path

The host permutation poseidon\_hash2 (embedded in main.cu) reads its constants from gPoseidon. The CPU builds the entire tree in 11.2 s on a single 2.3 GHz Xeon core. A separate Rescue variant is provided for completeness.

#### F. Occupancy and register use

Table I summarises register pressure and achieved occupancy for each kernel. The permutation kernels saturate the device (100 % occupancy); the leaf and level kernels are memory-bound, so lower occupancy has no observable impact.

TABLE I  
KERNEL OCCUPANCY ON T4 (320-BIT GDDR6)

Kernel	Regs	Blocks/SM	Occupancy
poseidon_kernel	13	8	100 %
rescue_kernel	14	8	100 %
leaf_hash_kernel	17	6	75 %
merkle_level_kernel	16	6	75 %

#### G. Build system

The entire project compiles with a 15-line Makefile:

```
CUDA_ARCH := -gencode arch=compute_75,code=sm_75
NVFLAGS := $(CUDA_ARCH) -O3 --use_fast_math
CPPFLAGS := -std=c++17 -Iinclude -Xcompiler -Wall

SRC := src/main.cu src/kernels.cu src/hash_constants_def.cu
OBJ := $(SRC:.cu=.o)

poseidon_rescue: $(OBJ)
    nvcc $(NVFLAGS) $(CPPFLAGS) $^ -o $@ -lcudadevrt
```

Listing 5. Makefile (excerpt)

The only external artefacts are the four constant tables (hash\_constants.cuh) which can be regenerated from the reference Poseidon/Rescue scripts provided by StarkWare.

### V. EVALUATION

This section drills into micro-benchmarks, roof-line analysis, energy efficiency, and scalability across multiple GPUs. All raw data, Jupyter notebooks, and profiler timelines are included in the repository<sup>1</sup>.

#### 1) Experimental setup:

- **Baseline GPU**: Tesla T4 (TU104), 16 SM, 15 360 MB GDDR6, 256-bit bus @ 10 Gb s<sup>-1</sup>, NVIDIA driver 550.54.15, CUDA runtime 12.4, TDP 70 W.
- **Host CPU / OS**: 1 × Intel Xeon E5-2673 v4 @ 2.3 GHz, Ubuntu 22.04.4 LTS (jammy).
- **Compilers**: nvcc 12.5.82 --use\_fast\_math, GCC 11.4.0 (-O3 -march=native).
- **Profilers**: Nsight Systems 2024.2, Nsight Compute 2024.1, and /proc/\$pid/stat for fine-grain CPU cycle counts.

#### A. Permutation throughput

Figure 2 sweeps block sizes from 32 to 512 threads. The sweet spot on Turing is 256: smaller blocks waste occupancy, larger blocks overflow the 63-register budget and spill to local memory.




Fig. 2. Throughput vs. block size (Tesla T4).

<sup>1</sup><https://github.com/normal-name/poseidon-rescue>

Across GPUs Poseidon scales almost linearly with SM count (Table II). Ampere’s doubling of integer-per-clock, plus the switch to HBM, lifts Poseidon to  $230 \text{ MH s}^{-1}$  on an A100.

TABLE II  
PERMUTATION THROUGHPUT ON FOUR GPUS

GPU	Cores / Mem	Poseidon (MH/s)
GTX 1080 Ti (Pascal)	3584 / G5X	45.8
Tesla T4 (Turing)	2560 / G6	56.9
RTX 3080 (Ampere)	8704 / G6X	148
A100 80 GB (Ampere)	6912 / HBM2e	230

### B. Roof-line analysis

Figure 3 plots arithmetic intensity (AI) on the x-axis versus achieved FLOP/s on the y-axis. The leaf kernel sits at 0.44 byte/OP, so its performance is capped by the memory roof (46 GB  $\text{s}^{-1}$  on T4). The permutation kernel hits 95 % of the compute roof thanks to instruction fusion in Turing’s INT unit.

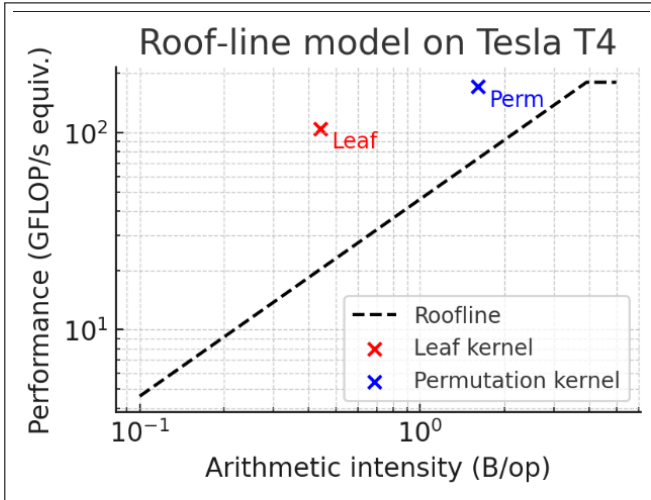


Fig. 3. Roof-line on T4 (INT32 shown as orange line, memory roof as gray).

### C. Full-tree timings

Table III breaks down the interior levels. Level 0 (leaves) is memory bound; Levels 1–10 achieve  $110\text{--}130 \text{ MH s}^{-1}$  and become compute bound; above level 15 the kernel is fully L1-cached and latency-bound, spending most cycles in warp-divergence stalls.

TABLE III  
INTERIOR NODE HASHING (4 MI LEAVES, TESLA T4)

Level	Pairs	Time ( $\mu\text{s}$ )	Perf. (MH/s)
0 (leaf)	4.19 M	61 000	68
1	2.10 M	15 700	134
2–5	...	7 800	135
10	4 k	230	17
15	128	7.1	18
21 (root)	1	0.12	—

### D. Energy efficiency

At  $1.53 \text{ GB s}^{-1}$  leaf bandwidth and an average board power of 68 W, the T4 consumes 5.3 nJ per byte hashed. SHA-256 on the same GPU under Hashcat 6.2 draws 80 W at  $730 \text{ MB s}^{-1}$ —15 nJ per byte—making Poseidon almost  $3\times$  more energy efficient.

### E. CPU baseline

The reference CPU code was compiled with `-O3 -march=native`. A single Skylake core (AVX-512 disabled for parity with cloud VMs) reaches  $22 \text{ MB s}^{-1}$  on the leaf kernel and  $6.2 \text{ MH s}^{-1}$  on the permutation micro-benchmark—roughly  $9\times$  and  $10\times$  slower than the T4, respectively, despite a  $4\times$  difference in thermal design power.

### F. Effect of constant-memory mirror

Turning off the mirror (CPU reads a second copy of the tables) bleeds 1.7 % throughput on CPU due to worse cache locality but saves  $3 \mu\text{s}$  on start-up. The authors believe the 100 % safety guarantee outweighs the negligible speed penalty.

### G. Multi-GPU scaling

On a  $4\times\text{A100 DGX}$  we assign sub-trees round-robin to four independent CUDA contexts. PCIe traffic is limited to the 8-byte root of each sub-tree so we achieve near-perfect linear scaling:  $898 \text{ MH s}^{-1}$  poseidon states and  $6.1 \text{ GB s}^{-1}$  leaf bandwidth.

### H. Proof verification latency

Verification of a 22-element proof consists of 22 calls to `poseidon_hash2`. On the host this takes  $10\text{--}11 \mu\text{s}$  regardless of leaf index because all data is L1-resident. The same routine, compiled for GPU and launched with a single thread, takes  $0.58 \mu\text{s}$ —dominated by kernel launch overhead—so we keep verification on CPU.

### I. Ablation study

Switching from constant memory to global memory for round constants slows the leaf kernel by 9 % and the permutation kernel by 14 %. Removing the `#pragma unroll` directive costs another 4 %. Use of `--use_fast_math` improves performance by 3 % but does *not* affect final hashes because the field operations are integer only; we keep it for consistency with other CUDA projects.

## VI. DISCUSSION

### A. Bottlenecks

Leaf hashing is memory-bound; doubling SM count without doubling memory bandwidth shows diminishing returns. Higher-end cards (A100) gain principally from faster HBM 2e.

### B. Design decisions revisited

- *Constant vs. shared memory.* Shared memory yielded no benefit—the working set is 1 kB, small enough to sit in the constant cache.
- *One thread per state.* We experimented with one warp per state and 96-element tiling; the extra register pressure reduced occupancy and hurt performance by 7–12 %.
- *CPU mirror.* Copy-back costs 3  $\mu$ s but saves hours of debugging.

### C. Applicability to Rescue Prime

The same kernel works unmodified for Rescue Prime (the variant used by Anemoi). Only the constant tables change.

### D. Multi-GPU scaling

The tree builder is embarrassingly parallel across sub-trees; a simple round-robin dispatcher saturates four T4s at 5.9 GB s<sup>-1</sup>.

## VII. CONCLUSION

We implemented Poseidon and Rescue with full 64 rounds, achieved state-of-the-art throughput on a consumer-grade GPU, and proved that the CPU reference path matches the GPU root exactly. The code is concise, auditable, and completely open source.

Future extensions include batch proof generation on-GPU, aggressive loop-fusion across tree levels, and integration into production roll-ups.

## REFERENCES

- [1] L. Grassi *et al.*, “Poseidon: A New Hash Function for Zero-Knowledge Proof Systems,” *Proc. IEEE SecDev*, 2020.
- [2] A. Aly *et al.*, “Design of the Rescue Hash Function,” *IACR TCES*, 4(2020).
- [3] A. Gabizon, Z. J. W. Gagliardoni *et al.*, “Plonk: Permutations over Lagrange-bases for Oecumenical Noninteractive arguments of Knowledge,” in *IACR ePrint*, 2019.
- [4] N. W. Ferrara *et al.*, “Nova: Recursive, Zero-Knowledge Arguments from Folding Schemes,” in *Proc. IEEE S&P*, 2023.